
Semantic Runtime Protocol: A Minimal Study of Bounded Semantic Drift in Long-Horizon LLM Interaction

Anonymous Author(s)

Abstract

Long-horizon LLM systems frequently manage context through prompt accumulation, summarization memory, or retrieval-based memory, but these strategies do not explicitly treat semantic state as an object that can be compressed, recovered, validated, and updated over time. We present Semantic Runtime Protocol (SRP), a semantic-state runtime abstraction for studying finite-horizon stability under repeated compression-recovery cycles. The paper makes a narrow claim: explicit semantic-state management can improve bounded semantic drift behavior relative to lightweight baselines while keeping token cost competitive over a finite horizon. We formalize SRP as a simple closed-loop transformation protocol, use bounded semantic drift as the main analysis lens, and outline a reproducible experiment comparing SRP with lightweight baselines on long-term consistency tasks. The contribution of this semester-version paper is a small but testable framing for long-term LLM stability, not a full semantic operating system.

1 Introduction

Large language models (LLMs) have enabled a new paradigm of interactive and task-oriented computation. However, their ability to operate over long horizons remains fundamentally constrained by how contextual information is represented, maintained, and transformed across interactions.

Existing systems typically rely on three paradigms: prompt-based context engineering, retrieval-augmented memory systems, and agent-based architectures. While effective in short or medium horizons, these approaches lack a unified and formally grounded mechanism for managing semantic information as a persistent computational object.

As a result, long-horizon LLM systems suffer from three persistent limitations: (1) uncontrolled semantic degradation under compression or summarization, (2) lack of verifiable recovery of intermediate semantic states, and (3) absence of a principled lifecycle model for evolving semantic representations over time.

We argue that the fundamental limitation is not the absence of memory or retrieval mechanisms, but the lack of a runtime abstraction for semantic state itself. In current systems, context is treated either as unstructured token sequences, externalized but loosely structured documents, or implicit hidden state within agent trajectories. None of these representations cleanly supports explicit transformation semantics, meaning that compression, recovery, and update operations are not governed by a verifiable state transition model.

In this work, we propose that long-horizon LLM systems should be modeled not simply as memory-augmented agents, but as a semantic state runtime. We introduce SRP, a framework that treats semantic information as a first-class computational object defined by structured state and explicit transformation operators. The key distinction is that memory stores state, while runtime governs state transition, validation, recovery, and update.




Figure 1: SRP as a semantic-state runtime layer between prompt/context management and downstream LLM reasoning.

Under SRP, semantic state is not static context, but a dynamic object that can be compressed into compact representations, recovered into operational semantic forms, validated through consistency checks, and updated through controlled transitions. This reframes long-horizon interaction as a closed-loop state transformation problem rather than an open-ended context retention problem.

The central insight of SRP is that long-horizon stability depends not on preserving raw context exactly, but on maintaining bounded semantic equivalence under repeated transformation. We therefore introduce the notion of bounded semantic drift, which measures how far recovered semantic state deviates from original task-relevant behavior over time.

This paper makes three contributions:

1. We introduce SRP as a semantic-state runtime abstraction for long-horizon LLM interaction.
2. We define a minimal protocol of compression, recovery, validation, and update operators over semantic state.
3. We propose bounded semantic drift as a practical formal and empirical lens, and outline an evaluation protocol against prompt, summary, and retrieval baselines.

Our goal in this semester-version paper is deliberately narrow. We do not claim universal semantic correctness, full convergence, or replacement of all memory and agent systems. Instead, we present a bounded-scope runtime abstraction and a finite-horizon evaluation framework that can be tested empirically.

Failure Boundary

To keep the claim falsifiable, the paper should also state when SRP fails. SRP is not guaranteed to work if the semantic vocabulary is corrupted, the validator is unreliable, the recovery step becomes non-invertible in practice, or cumulative drift exceeds the tolerance budget. This failure boundary improves credibility because it turns SRP into a testable runtime claim rather than an unconditional promise.

2 Related Work

Research on long-horizon LLM systems has grown rapidly, but most existing approaches manage context indirectly through prompts, summaries, retrieval, or agent trajectories. SRP is related to these directions, but differs in its central claim: semantic state itself should be modeled as an explicit runtime object with governed transformation operators.

2.1 Prompt Compression and Context Reduction

One major line of work focuses on reducing prompt length while preserving task performance [Jiang et al., 2023, Liskavets et al., 2025, Pan et al., 2024]. These methods show that raw prompt text often contains redundancy and can be shortened substantially without immediate loss in downstream utility.

However, prompt compression methods generally optimize prompts as token sequences rather than formalizing semantic state as a persistent computational object. Compression is typically a preprocessing step over context, not a runtime protocol with explicit recovery, validation, and update.

2.2 Memory-Augmented LLM Systems

Another major direction augments LLMs with external memory, including summary-based memory, episodic buffers, vector databases, and retrieval-augmented generation [Lewis et al., 2020, Wang et al., 2023a,b, Wu et al., 2025, Shan et al., 2025]. These systems improve personalization and continuity by storing or retrieving prior information across turns.

Their limitation, for our purposes, is not lack of utility but lack of explicit transformation semantics. Summarization-based memory replaces earlier context with a lossy abstraction, while retrieval-based memory selects prior records without defining which semantic invariants should be preserved after reuse.

2.3 Agent Frameworks and Trajectory-Based Systems

Agent systems externalize long-horizon behavior through plans, tool calls, reasoning traces, and action trajectories. These systems are effective for structured task completion, but generally treat history as execution process rather than as a dedicated semantic state abstraction.

2.4 Semantic Communication and Semantic Calibration

SRP is also related to work on semantic communication, concept-level representation, and semantic calibration [Guo et al., 2024, Nguyen et al., 2025, Meng et al., 2025]. These literatures motivate the shift from token fidelity to semantic fidelity and support the view that meaning preservation matters more than exact surface reproduction.

2.5 Reproducibility, Replay, and LLM Evaluation

Work on reproducibility, replay, prompt sensitivity, and experiment logging motivates the measurement side of SRP [Lu et al., 2023, Razavi et al., 2025, Wang and Wang, 2025, Semmelrock et al., 2024, Chen et al., 2022, Staudinger et al., 2024]. Because semantic drift accumulates over multiple steps, replayable experiments and stable evaluation infrastructure are especially important.

3 Runtime State and Operators

We model the runtime state at interaction step t as:

$$S_t = (M_t, V_t, P_t) \tag{1}$$

where M_t denotes structured semantic memory, V_t denotes vocabulary state, and P_t denotes the policy state governing transformation.

Minimal SRP loop placeholder.

Figure 2: Minimal SRP loop used in the semester-version experiment.

SRP defines four operators over semantic state:

$$C : S_t \rightarrow Z_t \quad (2)$$

$$R : Z_t \rightarrow S'_t \quad (3)$$

$$\text{Val} : (S_t, S'_t) \rightarrow [0, 1] \quad (4)$$

$$U : (S'_t, F_t) \rightarrow S_{t+1} \quad (5)$$

Compression maps the current semantic state to a compact representation Z_t . Recovery reconstructs an operational state from that compact form. Validation estimates whether task-relevant meaning is preserved. Update revises the state based on recovered semantics and observed feedback.

4 Formalization

The purpose of the formalization is practical rather than philosophical. We do not attempt to prove semantic truth. Instead, we define measurable error over observable task behavior.

Let O be a task-relevant observation operator that maps a semantic state to an observable behavior space:

$$O : S \rightarrow Y \quad (6)$$

We define semantic error between original and recovered state as:

$$\epsilon(S, S') = d(O(S), O(S')) \quad (7)$$

The goal is bounded degradation:

$$\epsilon(S, S') \leq \tau \quad (8)$$

Consider repeated transformation:

$$S_0 \xrightarrow{C} Z_0 \xrightarrow{R} S_1 \xrightarrow{C} Z_1 \xrightarrow{R} S_2 \rightarrow \dots \quad (9)$$

We define cumulative drift at step k as:

$$\Delta_k = d(O(S_0), O(S_k)) \quad (10)$$

We say that an SRP instance satisfies bounded semantic drift over horizon T if there exists a tolerance sequence $\{\tau_k\}$ such that for all $k \leq T$:

$$\Delta_k \leq \tau_k \quad (11)$$

If each compression-recovery step incurs local error at most e_t , and the update operator magnifies incoming error by at most factor λ , then after k cycles:

$$\Delta_k \leq \sum_{i=0}^{k-1} \lambda^{k-1-i} e_i \quad (12)$$

If $e_i \leq e$ for all i , then:

$$\Delta_k \leq e \sum_{i=0}^{k-1} \lambda^i \quad (13)$$

This yields linear growth when $\lambda = 1$ and a finite geometric bound when $0 \leq \lambda < 1$.

Drift over repeated compression-recovery cycles placeholder.

Figure 3: Drift over repeated compression-recovery cycles.

Token cost versus performance placeholder.

Figure 4: Token cost versus performance for SRP and baselines.

4.1 Why Drift Is the Main Metric

Drift is the primary metric because it is the shared failure mode across memory, retrieval, compression, and agent-based long-horizon systems. If repeated transformation is the operational setting, then drift is the most direct way to measure whether the semantic runtime abstraction is preserving task-relevant meaning over time.

In that sense, drift is not just one metric among several. It is the unifying quantity that the paper uses to compare systems across different context-management strategies.

5 Experimental Setup

We evaluate SRP in a finite-horizon setting designed to isolate bounded semantic drift under repeated transformation while keeping the study small enough to remain reproducible and interpretable. The goal is not to claim broad benchmark coverage, but to test whether explicit semantic-state management provides measurable stability benefits relative to lightweight baselines.

5.1 Main Experimental Question

The central experimental question is:

Does SRP maintain more stable task-relevant semantics over repeated compression-recovery cycles than lightweight baselines?

5.2 Main Methods

The main comparison focuses on four methods: raw prompt, summarization, retrieval-based memory, and SRP.

5.3 Main Tasks

We use three task families: multi-turn instruction consistency, long-context summarization and regeneration, and iterative compression-recovery cycles.

5.4 Benchmark and Replay

To improve replicability, the paper should treat benchmark design and replayability as part of the contribution. A minimal benchmark is useful only if it can be replayed under the same prompt template, backend, model version, and evaluation settings. The repo therefore serves not only as implementation support, but also as a replayable record of the state transitions that the paper studies.

6 Results

The semester version is designed to support a small number of clear empirical observations rather than a large benchmark sweep. The results section therefore prioritizes interpretability, reproducibility, and alignment with the paper’s central claim.

6.1 Drift Over Repeated Cycles

The primary result is the semantic-drift-over-iterations analysis. This figure measures how quickly task-relevant semantic information degrades when each method is repeatedly subjected to compression-recovery cycles.

The main question is whether SRP exhibits a slower increase in cumulative semantic drift than lightweight baselines. If the SRP curve grows more slowly or remains flatter over the finite horizon, this would support the core claim that explicit semantic-state management can stabilize repeated transformation.

6.2 Quality and Task Preservation

The second result concerns whether reduced drift corresponds to better preservation of task-relevant behavior. This is measured through task success under the selected core task families.

6.3 Efficiency

The third result concerns token cost. Since long-horizon systems are often constrained by context cost, the paper should evaluate whether SRP remains competitive in token usage while attempting to preserve semantic stability.

6.4 Camera-Ready Comparison

For a concise paper-facing view, the project also includes a camera-ready SRP versus strongest-baseline table. Its role is not to replace the fuller quality and efficiency analysis, but to provide a compact comparison that makes the central tradeoff immediately visible.

Under the current anchor-guided tuned local pilot, the clearest positive signal is that SRP now looks more balanced than plain summarization at ‘5’ and ‘7’ cycles, with higher task success and lower drift, even though retrieval-based memory still remains the strongest drift baseline on the present toy task suite.

7 Discussion

The main conceptual point of SRP is not that memory, retrieval, or agents are unnecessary. Rather, these components operate at adjacent layers. SRP targets the semantic-state layer itself.

8 Limitations

This work has several limitations. Drift is measured through proxies rather than direct semantic ground truth, results may depend on prompt design and judge model choice, and the current implementation uses a simplified version of SRP rather than the full blueprint.

9 Open Questions To Preserve

These should remain visible in the paper rather than being forced into a premature closure:

- Is SRP primarily a runtime abstraction, a protocol, or a semantic communication layer?
- What is the minimal semantic-state object?
- Which semantic drift metric should be considered the main one in the first paper?
- What failure cases should be shown explicitly?

10 Future Work

Several extensions are natural but should remain outside the first paper’s main claim: richer semantic state schemas, stronger validator design, cross-model and cross-domain generalization, semantic package composition and conflict resolution, and replay databases.

11 Conclusion

We introduced SRP as a semantic-state runtime abstraction for long-horizon LLM interaction. The main takeaway is modest but important: long-horizon LLM stability may benefit from explicit semantic runtime design, and this hypothesis is concrete enough to formalize, implement, and test. The long-term vision remains broader, but the present paper is designed to answer one falsifiable question about this framework rather than many future-facing ones.

Acknowledgments

This section is intentionally left blank for the draft. Replace it in the camera-ready version.

A Additional Details

The appendix can hold expanded proofs, additional experimental settings, ablations, or extra qualitative examples.

References

- Boyuan Chen, Mingzhi Wen, Yong Shi, Dayi Lin, Gopi Krishnan Rajbahadur, and Zhen Ming Jiang. Towards training reproducible deep learning models. *arXiv preprint arXiv:2202.02326*, 2022. URL <https://arxiv.org/pdf/2202.02326>.
- Shaolong Guo, Yuntao Wang, Ning Zhang, Zhou Su, Tom H. Luan, Zhiyi Tian, and Xuemin Shen. A survey on semantic communication networks: Architecture, security, and privacy. *arXiv preprint arXiv:2405.01221*, 2024. URL <https://arxiv.org/abs/2405.01221>.
- Huiqiang Jiang, Qianhui Wu, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. LlmLingua: Compressing prompts for accelerated inference of large language models. *arXiv preprint arXiv:2310.05736*, 2023. URL <https://arxiv.org/abs/2310.05736>.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich K ttler, Mike Lewis, Wen-tau Yih, Tim Rockt schel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, 2020. URL <https://arxiv.org/abs/2005.11401>.
- Barys Liskavets, Maxim Ushakov, Shuvendu Roy, Mark Klibanov, Ali Etemad, and Shane Luke. Prompt compression with context-aware sentence encoding for fast and improved llm inference. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025. URL <https://ojs.aaai.org/index.php/AAAI/article/view/34639/36794>.
- Sheng Lu, Hendrik Schuff, and Iryna Gurevych. How are prompts different in terms of sensitivity? *arXiv preprint arXiv:2311.07230*, 2023. URL <https://arxiv.org/abs/2311.07230>.
- Rui Meng, Song Gao, Dayu Fan, Haixiao Gao, Yining Wang, Xiaodong Xu, Bizhu Wang, Suyu Lv, Zhidi Zhang, Mengying Sun, Shujun Han, Chen Dong, Xiaofeng Tao, and Ping Zhang. A survey of secure semantic communications. *arXiv preprint arXiv:2501.00842*, 2025. URL <https://arxiv.org/abs/2501.00842>.
- Loc X. Nguyen, Avi Deb Raha, Pyae Sone Aung, Dusit Niyato, Zhu Han, and Choong Seon Hong. A contemporary survey on semantic communications: Theory of mind-based semantic communication, generative ai-driven semantic communication, and deep joint source-channel coding. *arXiv preprint arXiv:2502.16468*, 2025. URL <https://arxiv.org/abs/2502.16468>.

- Zhuoshi Pan, Qianhui Wu, Huiqiang Jiang, Menglin Xia, Xufang Luo, Jue Zhang, Qingwei Lin, Victor Rühle, Yuqing Yang, Chin-Yew Lin, H. Vicky Zhao, Lili Qiu, and Dongmei Zhang. LlmLingua-2: Data distillation for efficient and faithful prompt compression. *arXiv preprint arXiv:2403.12968*, 2024. URL <https://arxiv.org/abs/2403.12968>.
- Amirhossein Razavi, Mina Soltangheis, Negar Arabzadeh, Sara Salamat, Morteza Zihayat, and Ebrahim Bagheri. Benchmarking prompt sensitivity in large language models. *arXiv preprint arXiv:2502.06065*, 2025. URL <https://arxiv.org/html/2502.06065v1>.
- Harald Semmelrock, Tony Ross-Hellauer, Simone Kopeinik, Dieter Theiler, Armin Haberl, Stefan Thalmann, and Dominik Kowald. Reproducibility in machine learning-based research. *arXiv preprint arXiv:2406.14325*, 2024. URL <https://arxiv.org/abs/2406.14325>.
- Lianlei Shan, Shixian Luo, Zezhou Zhu, Yu Yuan, and Yong Wu. Cognitive memory in large language models. *arXiv preprint arXiv:2504.02441*, 2025. URL <https://arxiv.org/abs/2504.02441>.
- Moritz Staudinger, Wojciech Kusa, Florina Piroi, Aldo Lipani, and Allan Hanbury. A reproducibility and generalizability study of large language model-based boolean query generation for systematic reviews. *arXiv preprint arXiv:2411.14914*, 2024. URL <https://arxiv.org/abs/2411.14914>.
- Julian Junyan Wang and Victor Xiaoqi Wang. Assessing consistency and reproducibility in the outputs of large language models in finance and accounting. *arXiv preprint arXiv:2503.16974*, 2025. URL <https://arxiv.org/abs/2503.16974>.
- Qingyue Wang, Yanhe Fu, Yanan Cao, Shuai Wang, Zhiliang Tian, and Liang Ding. Recursively summarizing enables long-term dialogue memory in large language models. *arXiv preprint arXiv:2308.15022*, 2023a. URL <https://arxiv.org/abs/2308.15022>.
- Weizhi Wang, Li Dong, Hao Cheng, Xiaodong Liu, Xifeng Yan, Jianfeng Gao, and Furu Wei. Augmenting language models with long-term memory. *arXiv preprint arXiv:2306.07174*, 2023b. URL <https://arxiv.org/abs/2306.07174>.
- Yaxiong Wu, Sheng Liang, Chen Zhang, Yichao Wang, Yongyue Zhang, Huifeng Guo, Ruiming Tang, and Yong Liu. A survey on memory mechanisms in the era of llms. *arXiv preprint arXiv:2504.15965*, 2025. URL <https://arxiv.org/abs/2504.15965>.